Introduction to performance and the memory hierarchy

Computer Systems
Oct. 1, 2018

Michael Kirkedal Thomsen

Based on slides by Randal E. Bryant and David R. O'Hallaron

Performance!?

- How do you define performance?
- What influences performance of a computer?

- 5 minutes discussion with friend

Performance Realities

◼ ***There's more to performance than asymptotic complexity***

◼ **Constant factors matter too!**

- Easily see 10:1 performance range depending on how code is written
- Must optimize at multiple levels:
  - algorithm, data representations, procedures, and loops

◼ **Must understand system to optimize performance**

- How programs are compiled and executed
- How modern processors + memory systems operate
- How to measure program performance and identify bottlenecks
- How to improve performance without destroying code modularity and generality

Today

- **Performance**
- **Storage technologies and trends**
- **Locality of reference**
- **Caching in the memory hierarchy**

Optimizing Compilers

- **Provide efficient mapping of program to machine**
  - register allocation
  - code selection and ordering (scheduling)
  - dead code elimination
  - eliminating minor inefficiencies

- **Don't (usually) improve asymptotic efficiency**
  - up to programmer to select best overall algorithm
  - big-O savings are (often) more important than constant factors
    - but constant factors also matter

- **Have difficulty overcoming "optimization blockers"**
  - potential memory aliasing
  - potential procedure side-effects

Limitations of Optimizing Compilers

- **Operate under fundamental constraint**
  - Must not cause any change in program behavior
    - Except, possibly when program making use of nonstandard language features
  - Often prevents it from making optimizations that would only affect behavior under pathological conditions.
- **Behavior that may be obvious to the programmer can be obfuscated by languages and coding styles**
  - e.g., Data ranges may be more limited than variable types suggest
- **Most analysis is performed only within procedures**
  - Whole-program analysis is too expensive in most cases
  - Newer versions of GCC do interprocedural analysis within individual files
    - But, not between code in different files
- **Most analysis is based only on *static* information**
  - Compiler has difficulty anticipating run-time inputs
- **When in doubt, the compiler must be conservative**

Memory Matters

```c
/* Sum rows is of n X n matrix a
   and store in vector b  */
void sum_rows1(double *a, double *b, long n) {
    long i, j;
    for (i = 0; i < n; i++) {
        b[i] = 0;
        for (j = 0; j < n; j++)
            b[i] += a[i*n + j];
    }
}
```

```
# sum_rows1 inner loop
.L4:
        movsd    (%rsi,%rax,8), %xmm0        # FP load
        addsd    (%rdi), %xmm0               # FP add
        movsd    %xmm0, (%rsi,%rax,8)        # FP store
        addq     $8, %rdi
        cmpq     %rcx, %rdi
        jne      .L4
```

- Code updates `b[i]` on every iteration
- Why couldn't compiler optimize this away?

Memory Aliasing

```
/* Sum rows is of n X n matrix a
   and store in vector b  */
void sum_rows1(double *a, double *b, long n) {
    long i, j;
    for (i = 0; i < n; i++) {
        b[i] = 0;
        for (j = 0; j < n; j++)
            b[i] += a[i*n + j];
    }
}
```

```
double A[9] =
  { 0,    1,    2,
    4,    8,  16},
   32,   64, 128};

double B[3] = A+3;

sum_rows1(A, B, 3);
```

Value of B:

```
init:  [4, 8, 16]
```

```
i = 0: [3, 8, 16]
```

```
i = 1: [3, 22, 16]
```

```
i = 2: [3, 22, 224]
```

- Code updates `b[i]` on every iteration
- Must consider possibility that these updates will affect program behavior

Removing Aliasing

```
/* Sum rows is of n X n matrix a
   and store in vector b  */
void sum_rows2(double *a, double *b, long n) {
    long i, j;
    for (i = 0; i < n; i++) {
        double val = 0;
        for (j = 0; j < n; j++)
            val += a[i*n + j];
        b[i] = val;
    }
}
```

```
# sum_rows2 inner loop
.L10:
        addsd    (%rdi), %xmm0        # FP load + add
        addq     $8, %rdi
        cmpq     %rax, %rdi
        jne      .L10
```

- No need to store intermediate results

Cycles Per Element (CPE)

■ **Convenient way to express performance of program that operates on vectors or lists**

■ **Length = n**

■ **In our case: CPE = cycles per OP**

■ **T = CPE*n + Overhead**

- CPE is slope of line

I/O Bus

CPU chip

Register file

ALU

System bus

Memory bus

Bus interface

I/O bridge

Main memory

I/O bus

USB controller

Graphics adapter

Disk controller

Expansion slots for other devices such as network adapters.

Mouse   **Keyboard**

Monitor

Disk

Today

- **Performance**
- **Storage technologies and trends**
- **Locality of reference**
- **Caching in the memory hierarchy**

Random-Access Memory (RAM)

- **Key features**
  - RAM is traditionally packaged as a chip.
  - Basic storage unit is normally a cell (one bit per cell).
  - Multiple RAM chips form a memory.

- **RAM comes in two varieties:**
  - SRAM (Static RAM)
  - DRAM (Dynamic RAM)

Enhanced DRAMs

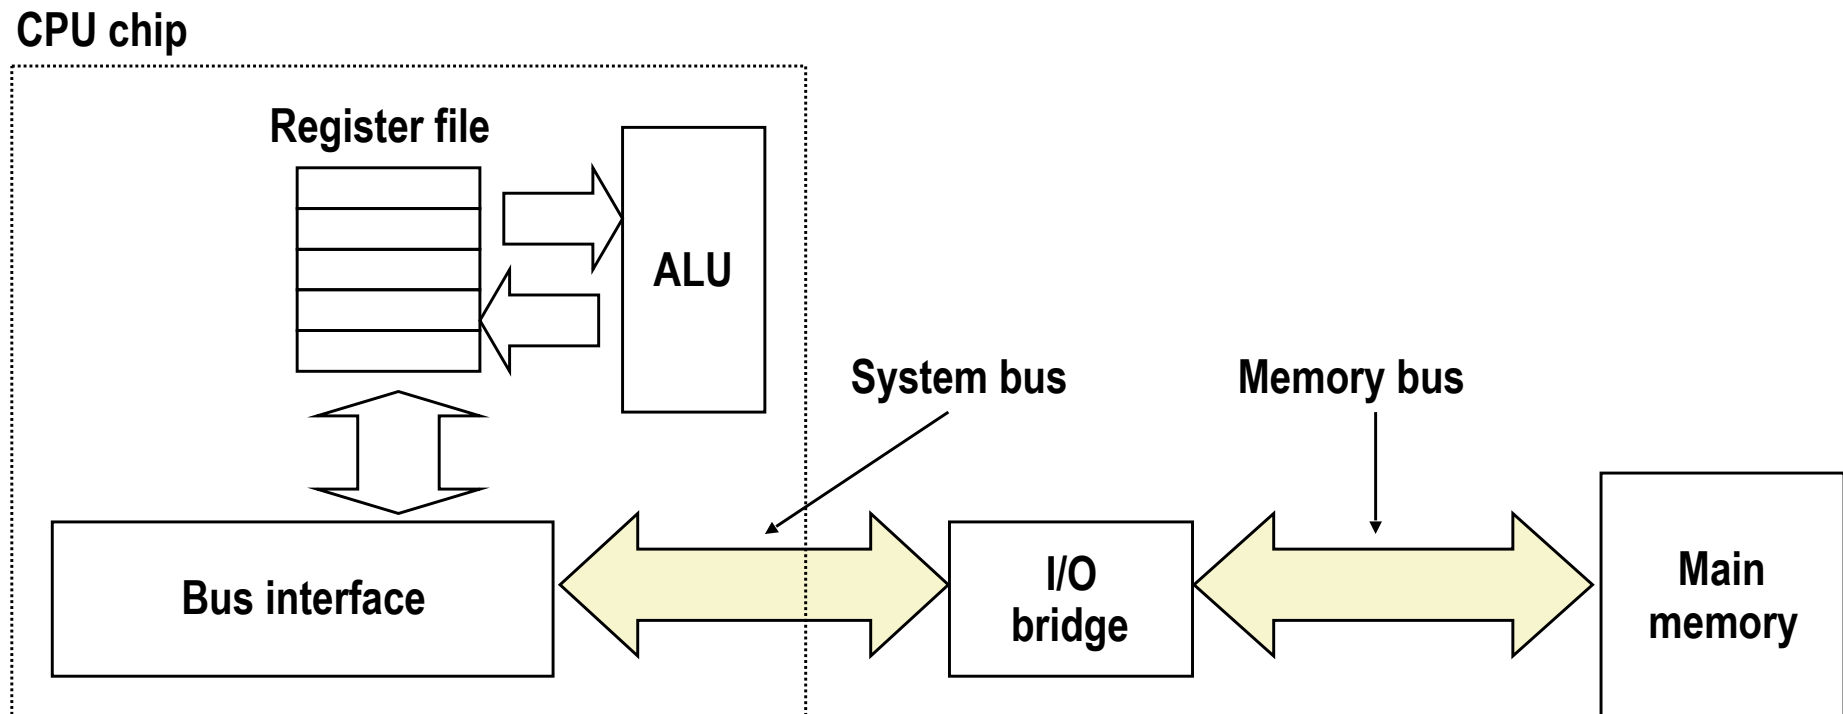- ◼ **Basic DRAM cell has not changed since its invention in 1966.**
    - ▪ Commercialized by Intel in 1970.
- ◼ **DRAM cores with better interface logic and faster I/O :**
    - ▪ Synchronous DRAM (SDRAM)
        - ▪ Uses a conventional clock signal instead of asynchronous control
        - ▪ Allows reuse of the row addresses (e.g., RAS, CAS, CAS, CAS)

    - ▪ Double data-rate synchronous DRAM (DDR SDRAM)
        - ▪ Double edge clocking sends two bits per cycle per pin
        - ▪ Different types distinguished by size of small prefetch buffer:
            - – DDR (2 bits), DDR2 (4 bits), DDR3 (8 bits)
        - ▪ By 2010, standard for most server and desktop systems
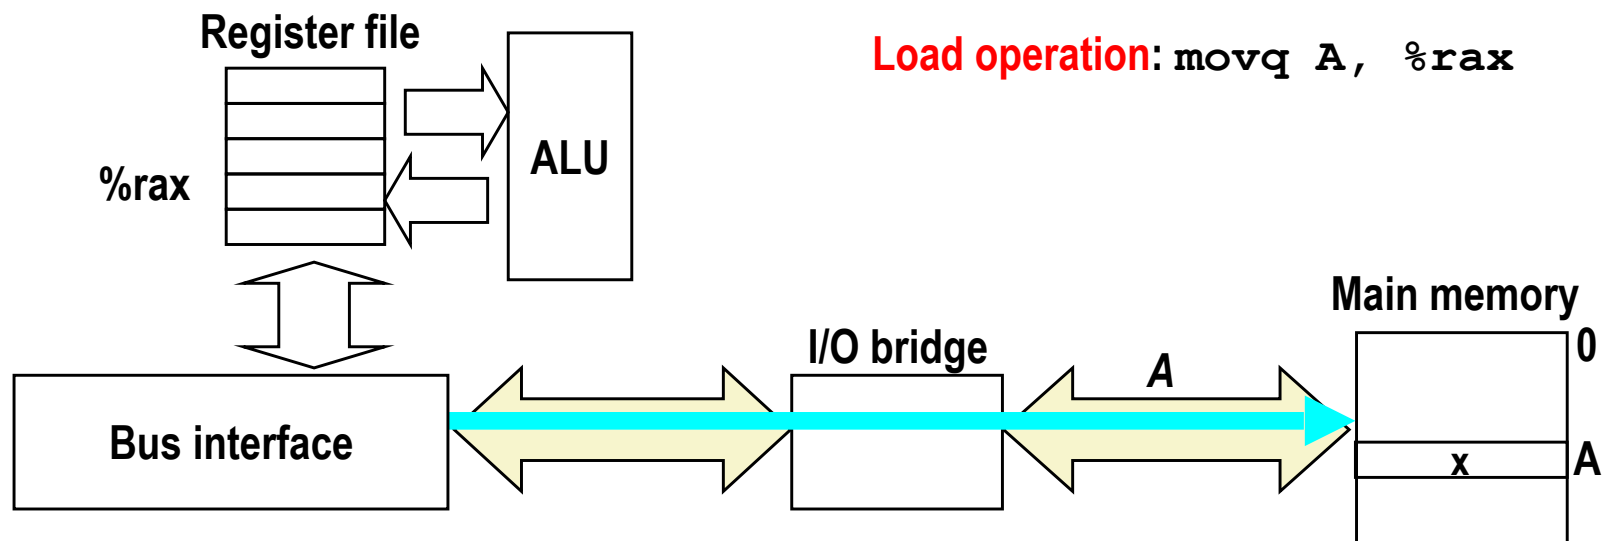   - ▪ Intel Core i7 supports only DDR3 SDRAM

Traditional Bus Structure Connecting CPU and Memory

- A **bus** is a collection of parallel wires that carry address, data, and control signals.

- Buses are typically shared by multiple devices.

CPU chip

Register file

ALU

System bus

Memory bus

Bus interface

I/O bridge

Main memory

Memory Read Transaction (1)

- ■ **CPU places address A on the memory bus.**

Register file

Load operation: `movq A, %rax`

ALU

%rax

Main memory

I/O bridge

A

0

Bus interface

x   **A**
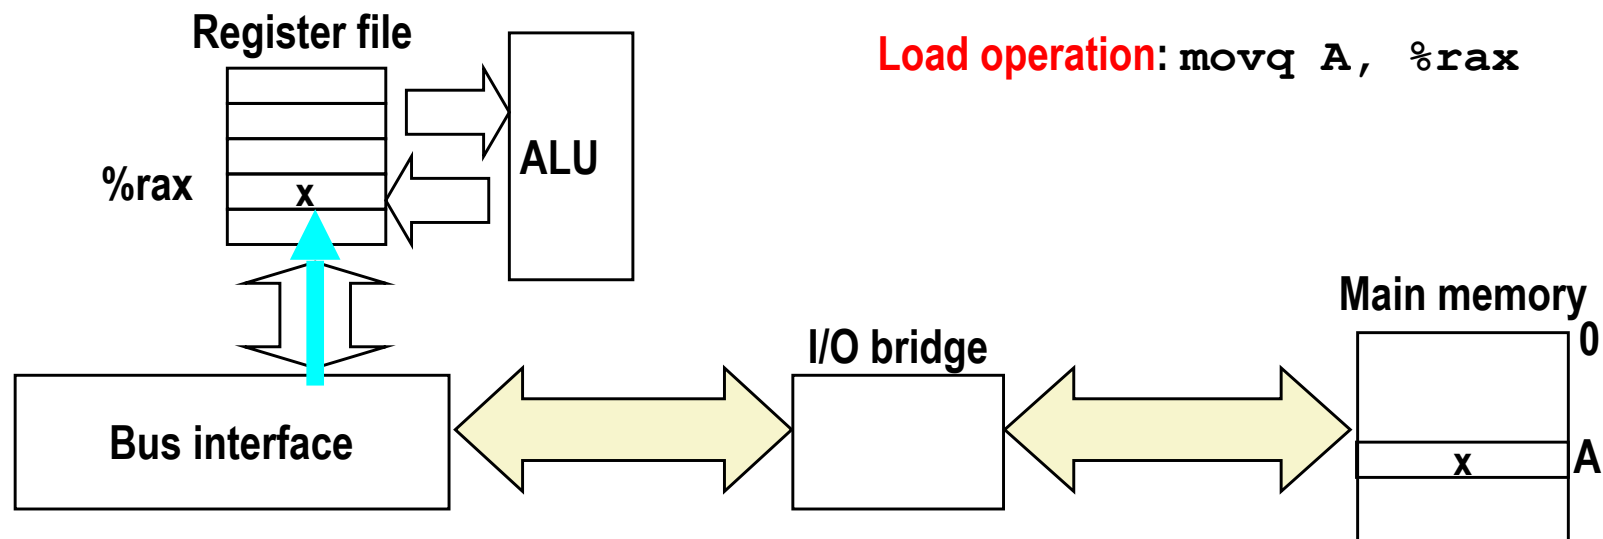
Memory Read Transaction (2)

- **Main memory reads A from the memory bus, retrieves word x, and places it on the bus.**

Register file

%rax

ALU

Load operation: `movq A, %rax`

Bus interface

I/O bridge

x

Main memory

0

x

A

Memory Read Transaction (3)

- CPU read word x from the bus and copies it into register %rax.

Register file

Load operation: `movq A, %rax`

ALU

%rax | x

Main memory

I/O bridge

Bus interface

0

x | A

Memory Write Transaction (1)

- CPU places address A on bus. Main memory reads it and waits for the corresponding data word to arrive.

Register file

Store operation: `movq %rax, A`

%rax   y

ALU

Bus interface

I/O bridge

A

Main memory

0

A

Memory Write Transaction (2)

- CPU places data word y on the bus.

Store operation: `movq %rax, A`

Register file

%rax — y

ALU

Bus interface

I/O bridge

y

Main memory

0

A

Memory Write Transaction (3)

- Main memory reads data word y from the bus and stores it at address A.

Store operation: `movq %rax, A`

Register file

%rax | y

ALU

main memory

Bus interface

I/O bridge

0

y | A

Conventional DRAM Organization

- **d x w DRAM:**
  - dw total bits organized as d supercells of size w bits

16 x 8 DRAM chip

Memory controller

(to/from CPU)

2 bits — addr

8 bits — data

cols
0 1 2 3

rows
0
1
2
3

supercell (2,1)

Internal row buffer
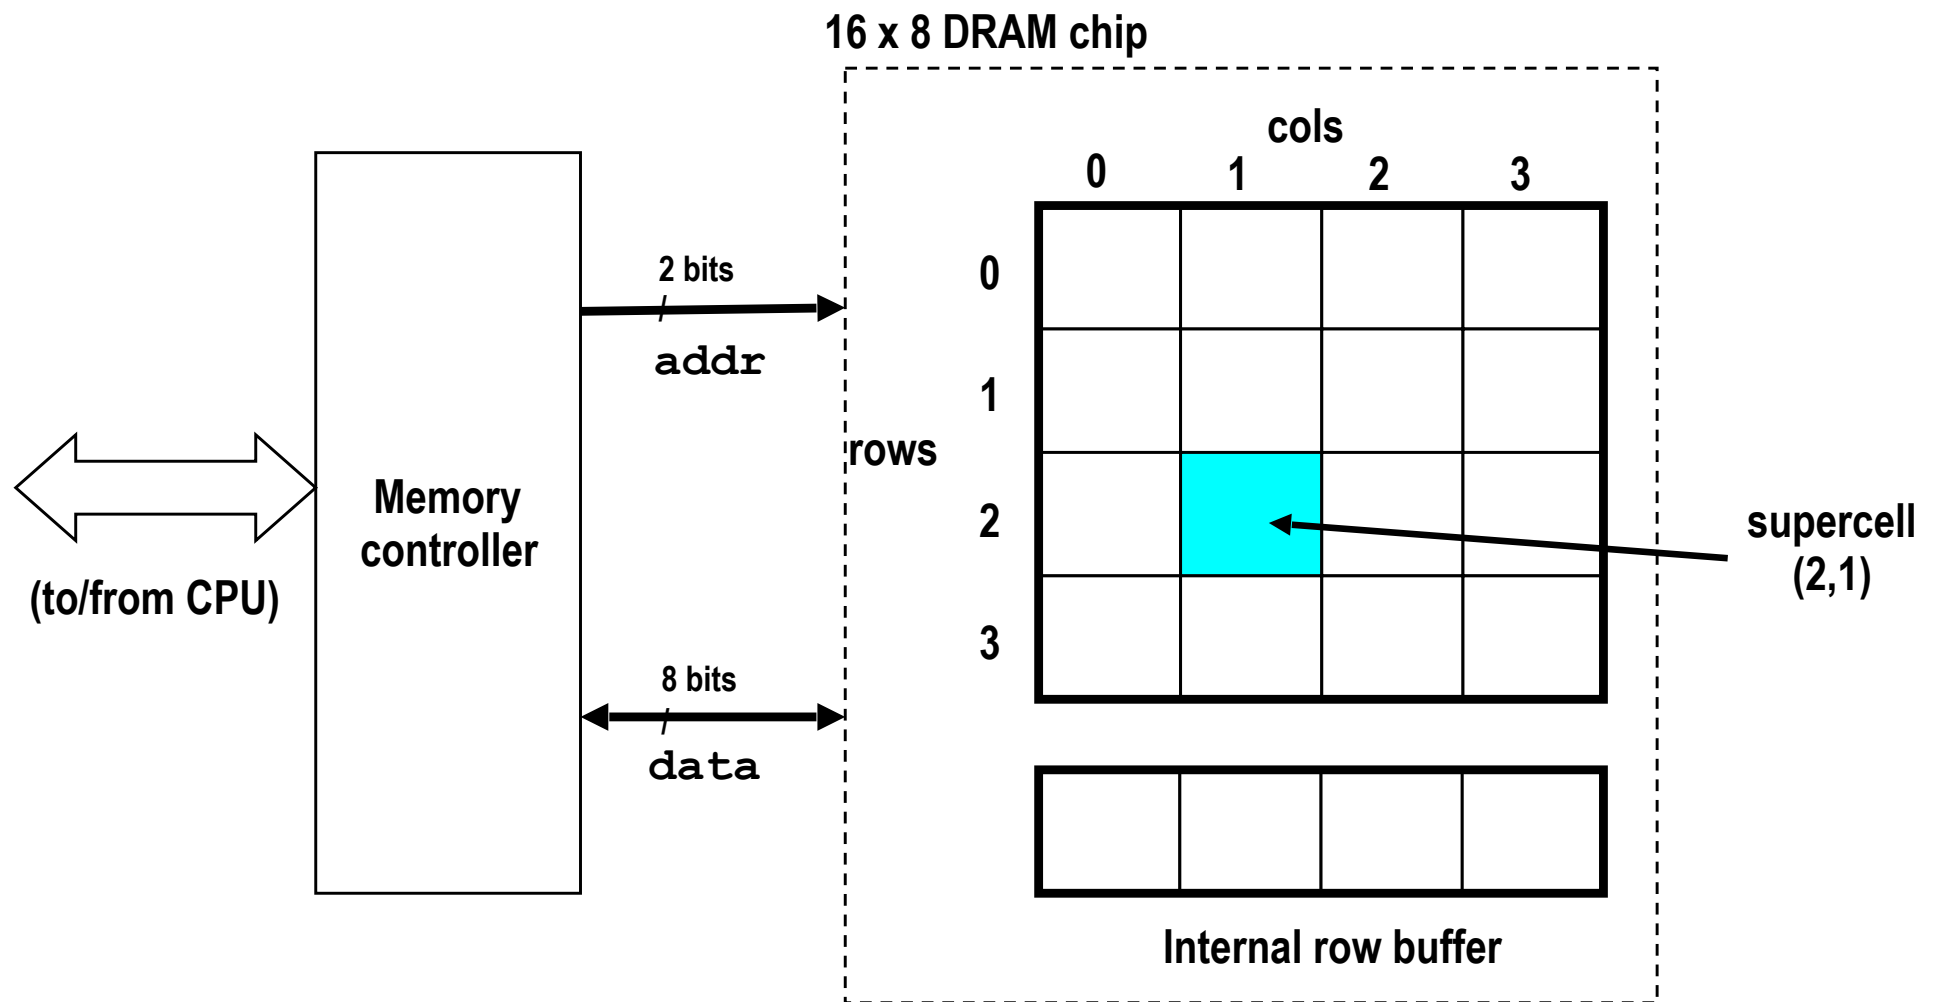
Reading DRAM Supercell (2,1)

Step 1(a): Row access strobe (RAS) selects row 2.

Step 1(b): Row 2 copied from DRAM array to row buffer.

16 x 8 DRAM chip

Cols

Memory controller

RAS = 2

addr

data

Rows

Internal row buffer
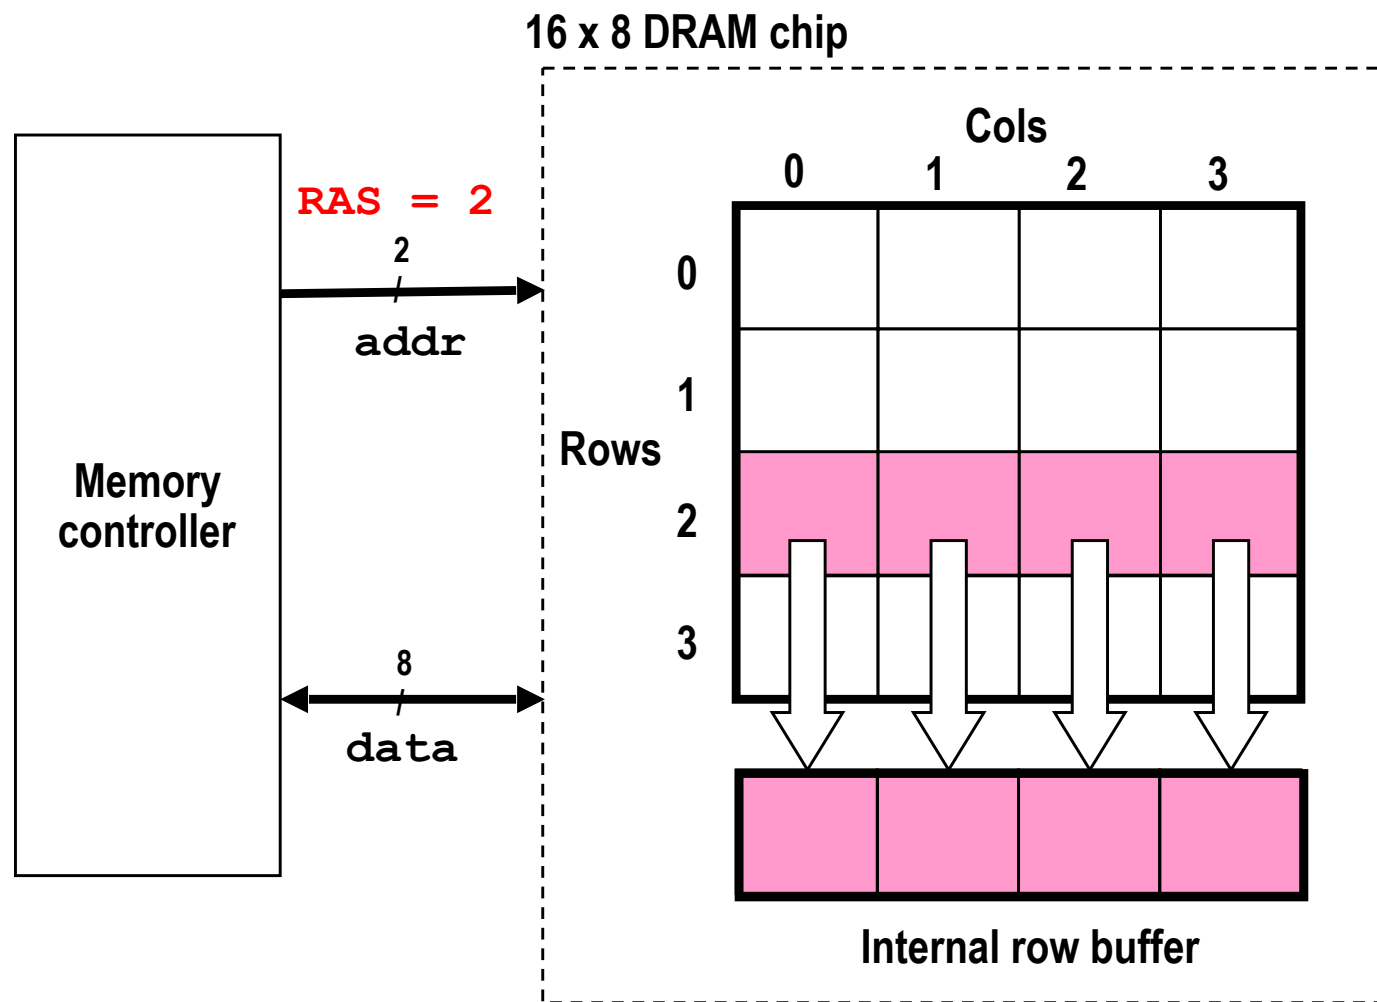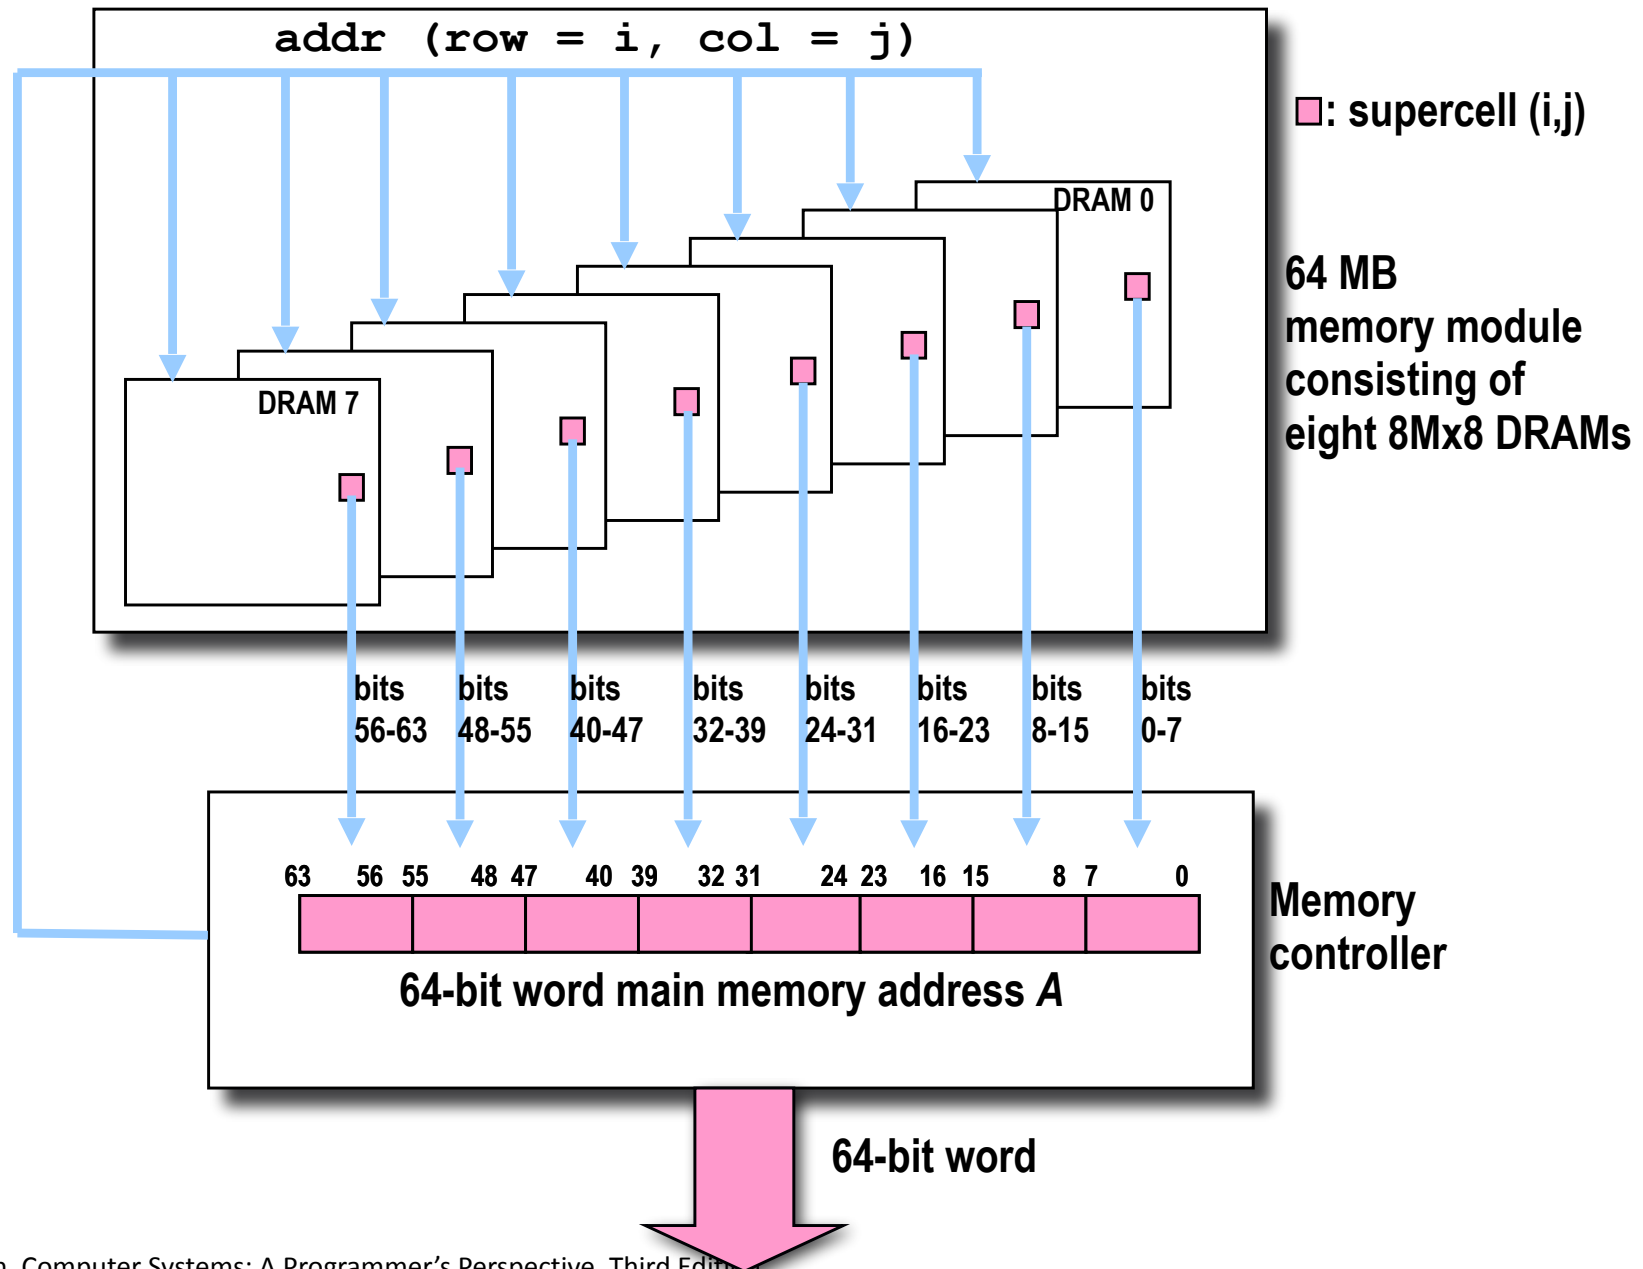
Reading DRAM Supercell (2,1)

Step 2(a): Column access strobe (CAS) selects column 1.

Step 2(b): Supercell (2,1) copied from buffer to data lines, and eventually back to the CPU.

Memory Modules

addr (row = i, col = j)

□: supercell (i,j)

DRAM 0

DRAM 7

64 MB
memory module
consisting of
eight 8Mx8 DRAMs

bits 56-63 | bits 48-55 | bits 40-47 | bits 32-39 | bits 24-31 | bits 16-23 | bits 8-15 | bits 0-7

63   56 55   48 47   40 39   32 31   24 23   16 15   8 7   0

Memory controller

64-bit word main memory address *A*

64-bit word
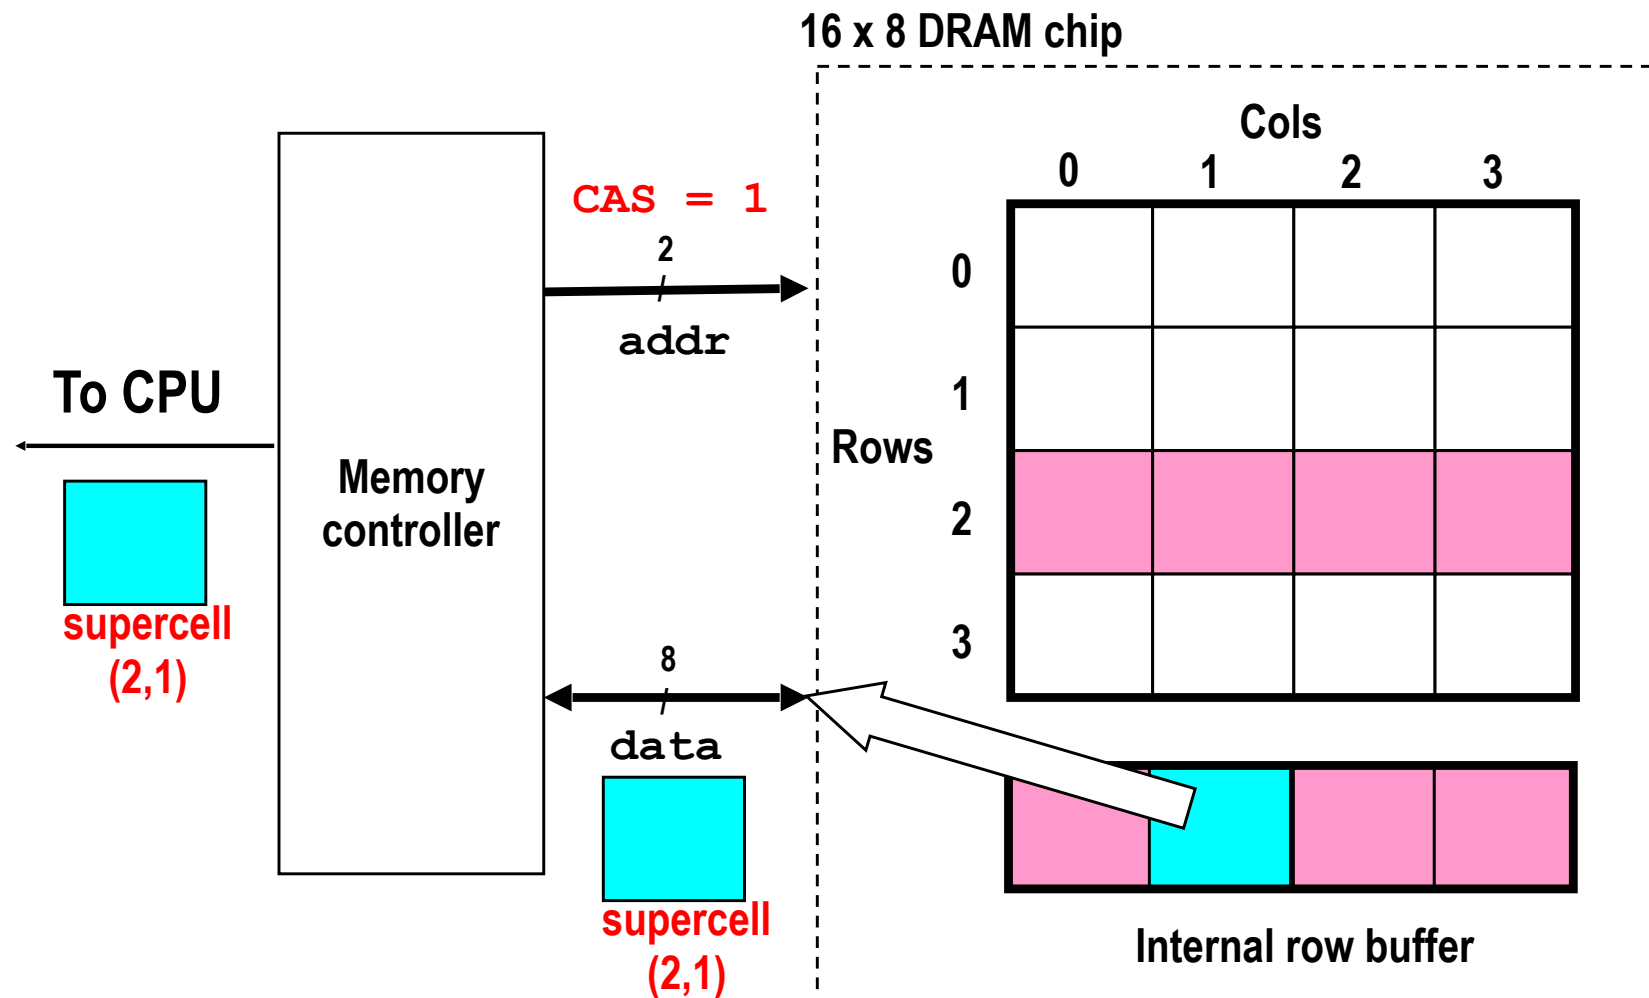
SRAM vs DRAM Summary

	Trans. per bit	Access time	Needs refresh?	Needs EDC?	Cost	Applications
SRAM	4 or 6	1X	No	Maybe	100x	Cache memories
DRAM	1	10X	Yes	Yes	1X	Main memories, frame buffers

Nonvolatile Memories

- **DRAM and SRAM are volatile memories**
  - Lose information if powered off.

- **Nonvolatile memories retain value even if powered off**
  - Read-only memory (ROM): programmed during production
  - Programmable ROM (PROM): can be programmed once
  - Eraseable PROM (EPROM): can be bulk erased (UV, X-Ray)
  - Electrically eraseable PROM (EEPROM): electronic erase capability
    - Flash memory: EEPROMs. with partial (block-level) erase capability
    - Wears out after about 100,000 erasings

- **Uses for Nonvolatile Memories**
  - Firmware programs stored in a ROM (BIOS, controllers for disks, network cards, graphics accelerators, security subsystems,…)
  - Solid state disks (replace rotating disks in thumb drives, smart phones, mp3 players, tablets, laptops,…)
 - Disk caches

What's Inside A Disk Drive?

Arm

Spindle

Platters

Actuator

Electronics
(including a
processor
and memory!)

SCSI
connector

Image courtesy of Seagate Technology

Disk Geometry

- **Disks consist of platters, each with two surfaces.**
- **Each surface consists of concentric rings called tracks.**
- **Each track consists of sectors separated by gaps.**

Disk Geometry (Muliple-Platter View)

■ **Aligned tracks form a cylinder.**

Disk Capacity

- **Capacity: maximum number of bits that can be stored.**
 - Vendors express capacity in units of gigabytes (GB), where 1 GB = 10^9 Bytes.
- **Capacity is determined by these technology factors:**
  - Recording density (bits/in): number of bits that can be squeezed into a 1 inch segment of a track.
  - Track density (tracks/in): number of tracks that can be squeezed into a 1 inch radial segment.
  - Areal density (bits/in2): product of recording and track density.

Recording zones

- **Modern disks partition tracks into disjoint subsets called recording zones**

    - Each track in a zone has the same number of sectors, determined by the circumference of innermost track.

    - Each zone has a different number of sectors/track, outer zones have more sectors/track than inner zones.

    - So we use **average** number of sectors/track when computing capacity.

Spindle

Computing Disk Capacity

Capacity = (# bytes/sector) x (avg. # sectors/track) x

(# tracks/surface) x (# surfaces/platter) x

(# platters/disk)

Example:

- 512 bytes/sector
- 300 sectors/track (on average)
- 20,000 tracks/surface
- 2 surfaces/platter
- 5 platters/disk

Capacity = 512 x 300 x 20000 x 2 x 5

= 30,720,000,000

= 30.72 GB

Disk Operation (Single-Platter View)

The disk surface spins at a fixed rotational rate

spindle

The read/write *head* is attached to the end of the *arm* and flies over the disk surface on a thin cushion of air.

By moving radially, the arm can position the read/write head over any track.

Disk Operation (Multi-Platter View)

Read/write heads
move in unison
from cylinder to cylinder

Arm

Spindle
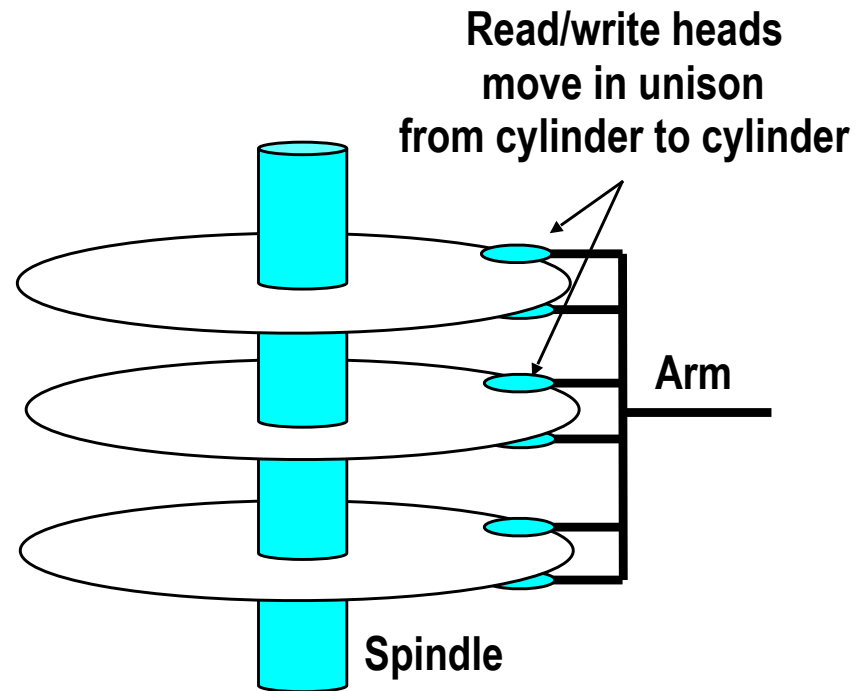
Disk Structure - top view of single platter

Surface organized into tracks

Tracks divided into sectors

Disk Access

Head in position above a track

Disk Access

Rotation is counter-clockwise

Disk Access – Read

About to read blue sector

Disk Access – Read

After **BLUE** read

After reading blue sector

Disk Access – Read

After BLUE read

Red request scheduled next

Disk Access – Seek

After **BLUE** read   Seek for **RED**

Seek to red's track

Disk Access – Rotational Latency

After **BLUE** read   Seek for **RED**   Rotational latency

Wait for red sector to rotate around

Disk Access – Read

After **BLUE** read   Seek for **RED**   Rotational latency   After **RED** read
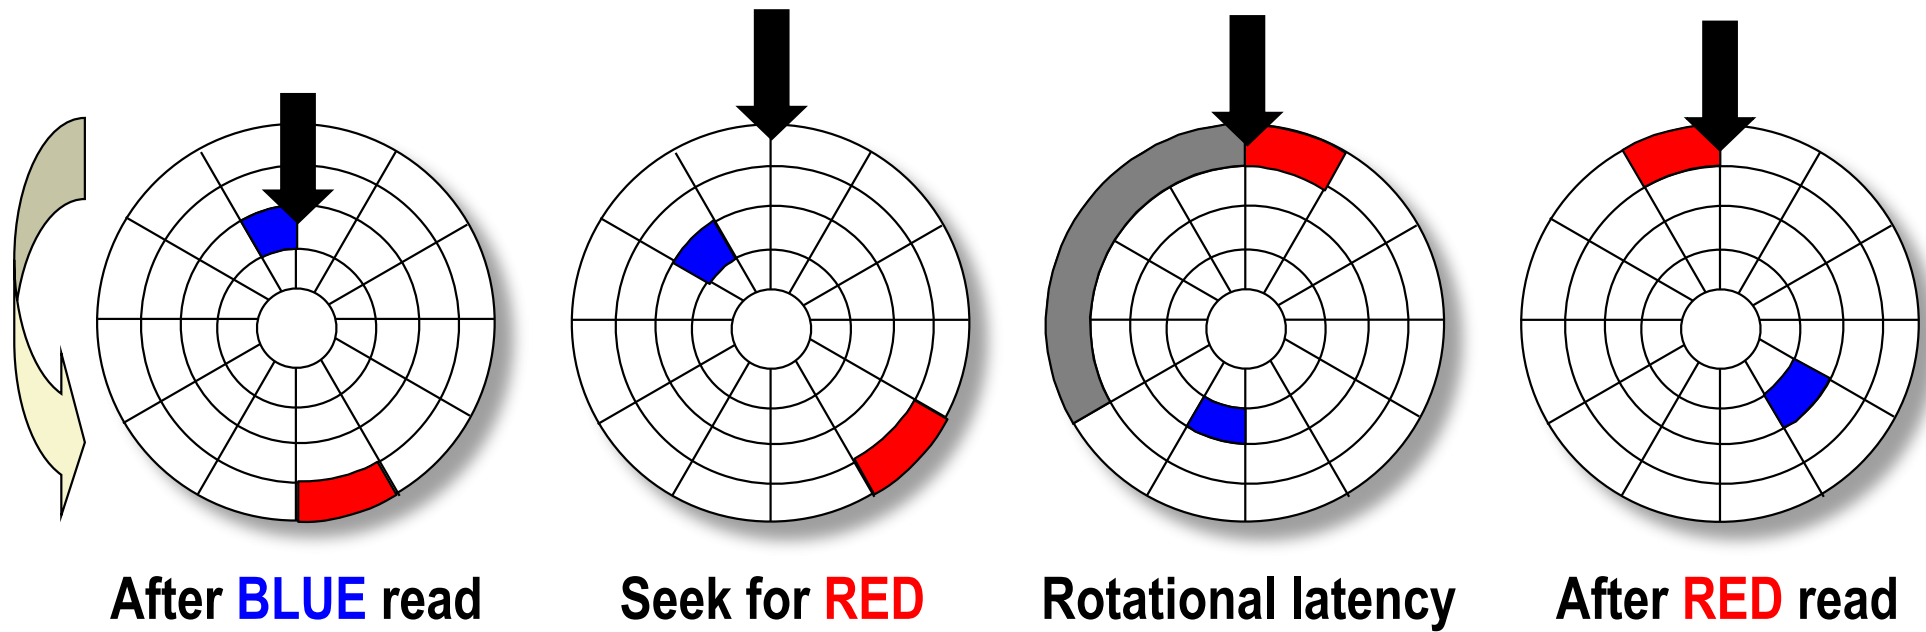
Complete read of red

Disk Access – Service Time Components

After BLUE read	Seek for RED	Rotational latency	After RED read
Data transfer	Seek	Rotational latency	Data transfer
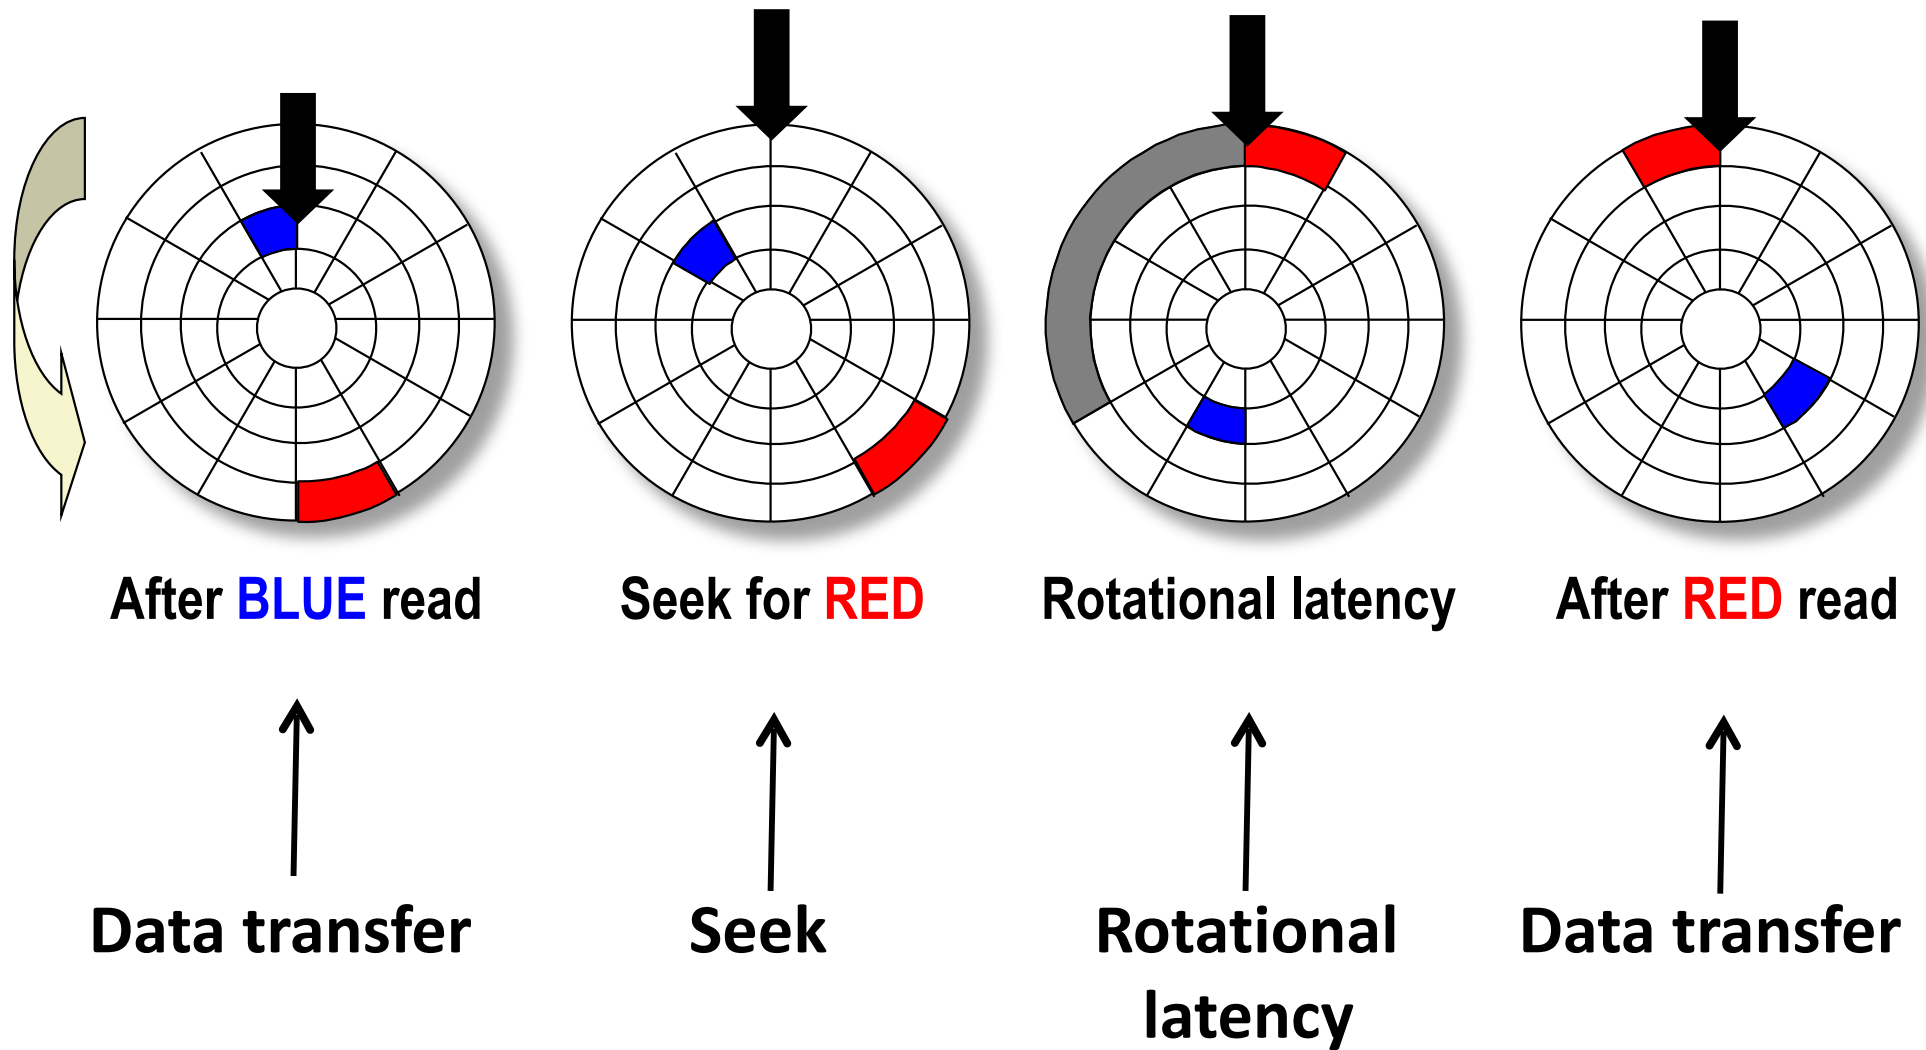
Disk Access Time

- **Average time to access some target sector approximated by :**
  - Taccess = Tavg seek + Tavg rotation + Tavg transfer

- **Seek time (Tavg seek)**
  - Time to position heads over cylinder containing target sector.
  - Typical Tavg seek is 3—9 ms

- **Rotational latency (Tavg rotation)**
  - Time waiting for first bit of target sector to pass under r/w head.
  - Tavg rotation = 1/2 x 1/RPMs x 60 sec/1 min
  - Typical Tavg rotation = 7200 RPMs

- **Transfer time (Tavg transfer)**
  - Time to read the bits in the target sector.
  - Tavg transfer = 1/RPM x 1/(avg # sectors/track) x 60 secs/1 min.

Disk Access Time Example

- **Given:**
  - Rotational rate = 7,200 RPM
  - Average seek time = 9 ms.
  - Avg # sectors/track = 400.

- **Derived:**
  - Tavg rotation = 1/2 x (60 secs/7200 RPM) x 1000 ms/sec = 4 ms.
  - Tavg transfer = 60/7200 RPM x 1/400 secs/track x 1000 ms/sec = 0.02 ms
  - Taccess = 9 ms + 4 ms + 0.02 ms

- **Important points:**
  - Access time dominated by seek time and rotational latency.
  - First bit in a sector is the most expensive, the rest are free.
  - SRAM access time is about 4 ns/doubleword, DRAM about 60 ns
    - Disk is about 40,000 times slower than SRAM,
    - 2,500 times slower than DRAM.

Logical Disk Blocks

- **Modern disks present a simpler abstract view of the complex sector geometry:**
  - The set of available sectors is modeled as a sequence of b-sized logical blocks (0, 1, 2, …)

- **Mapping between logical blocks and actual (physical) sectors**
  - Maintained by hardware/firmware device called disk controller.
  - Converts requests for logical blocks into (surface,track,sector) triples.

- **Allows controller to set aside spare cylinders for each zone.**
  - Accounts for the difference in "formatted capacity" and "maximum capacity".

I/O Bus

CPU chip

Register file

ALU

System bus

Memory bus

Bus interface

I/O bridge

Main memory

I/O bus

USB controller

Graphics adapter

Disk controller

Expansion slots for other devices such as network adapters.

Mouse   Keyboard

Monitor
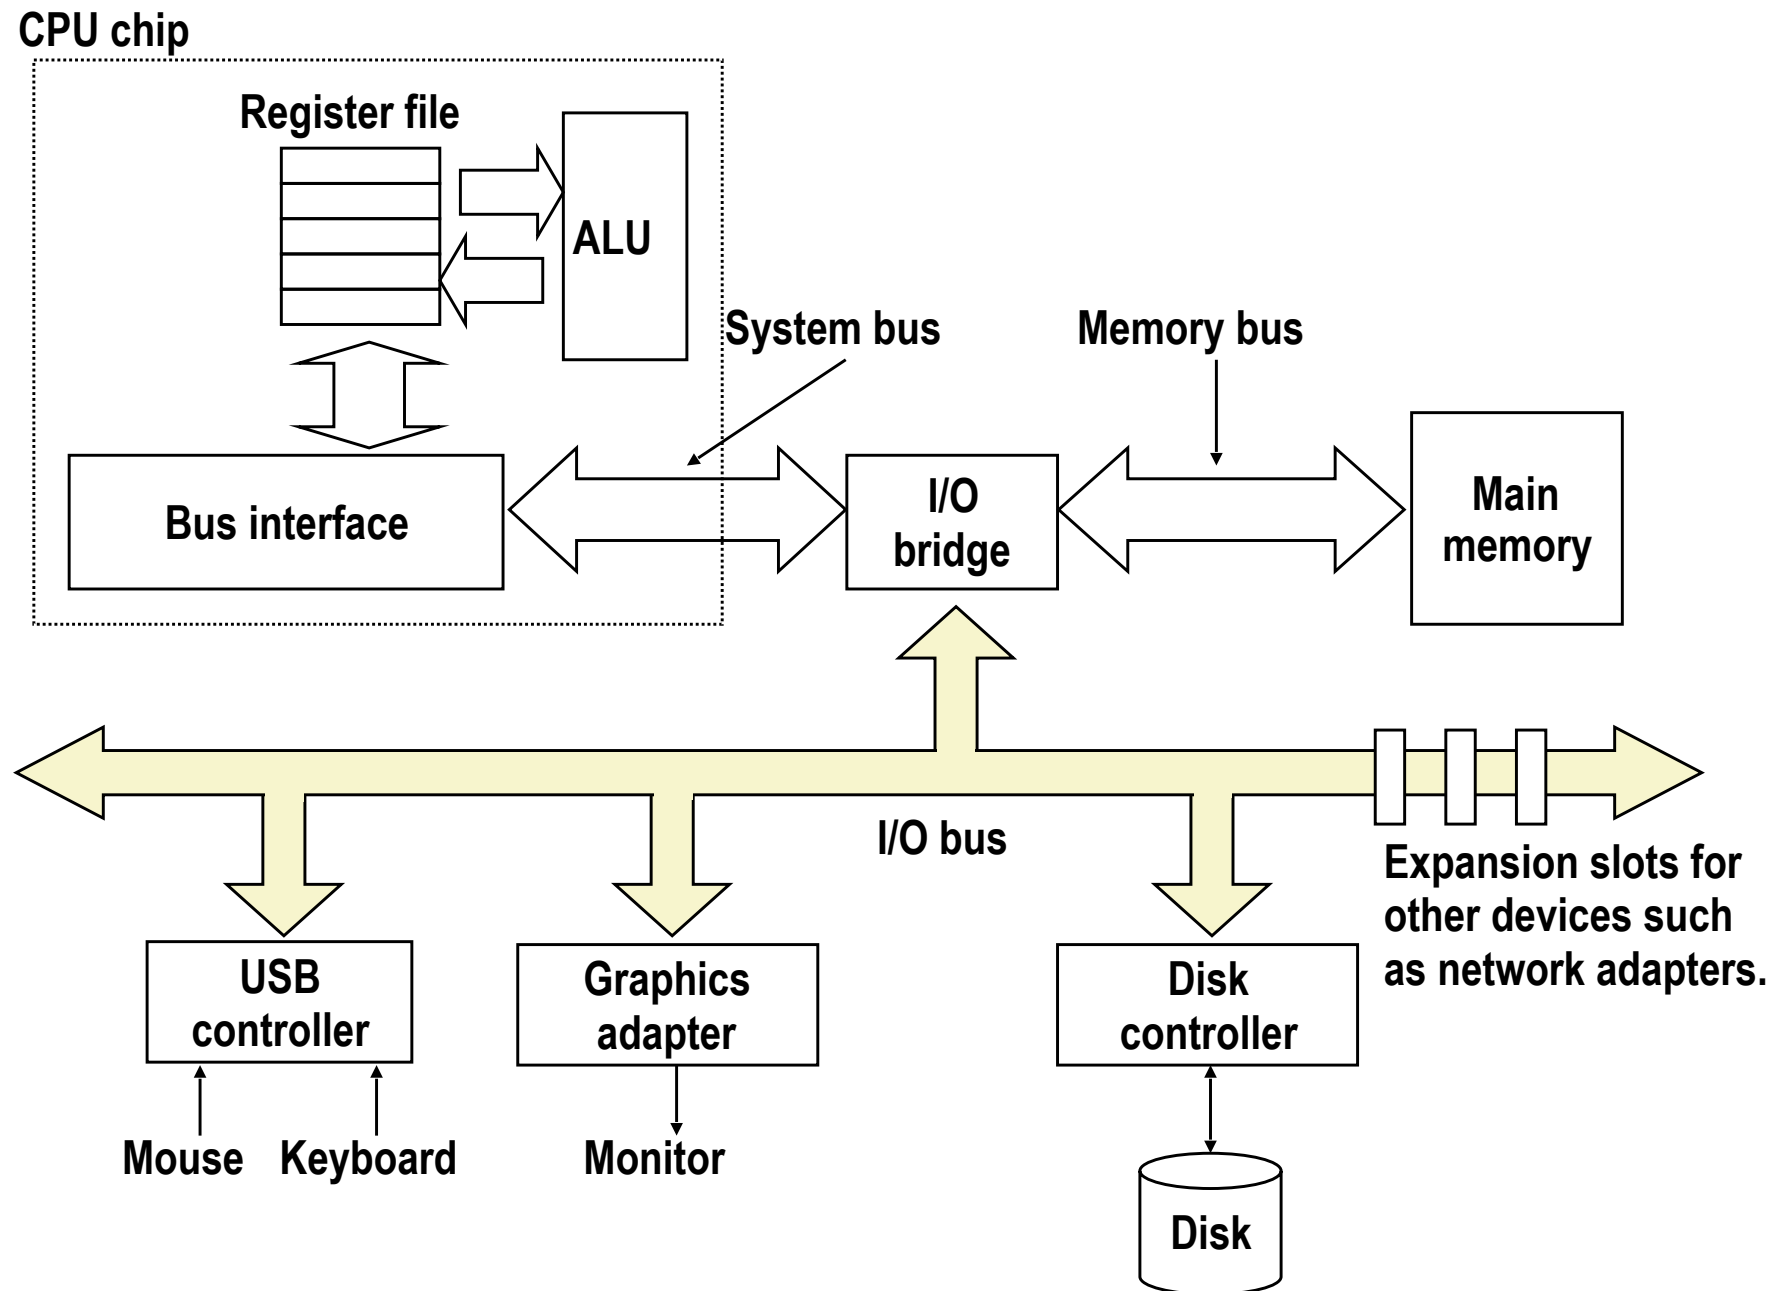
Disk

Reading a Disk Sector (1)

CPU chip

Register file

ALU

Bus interface
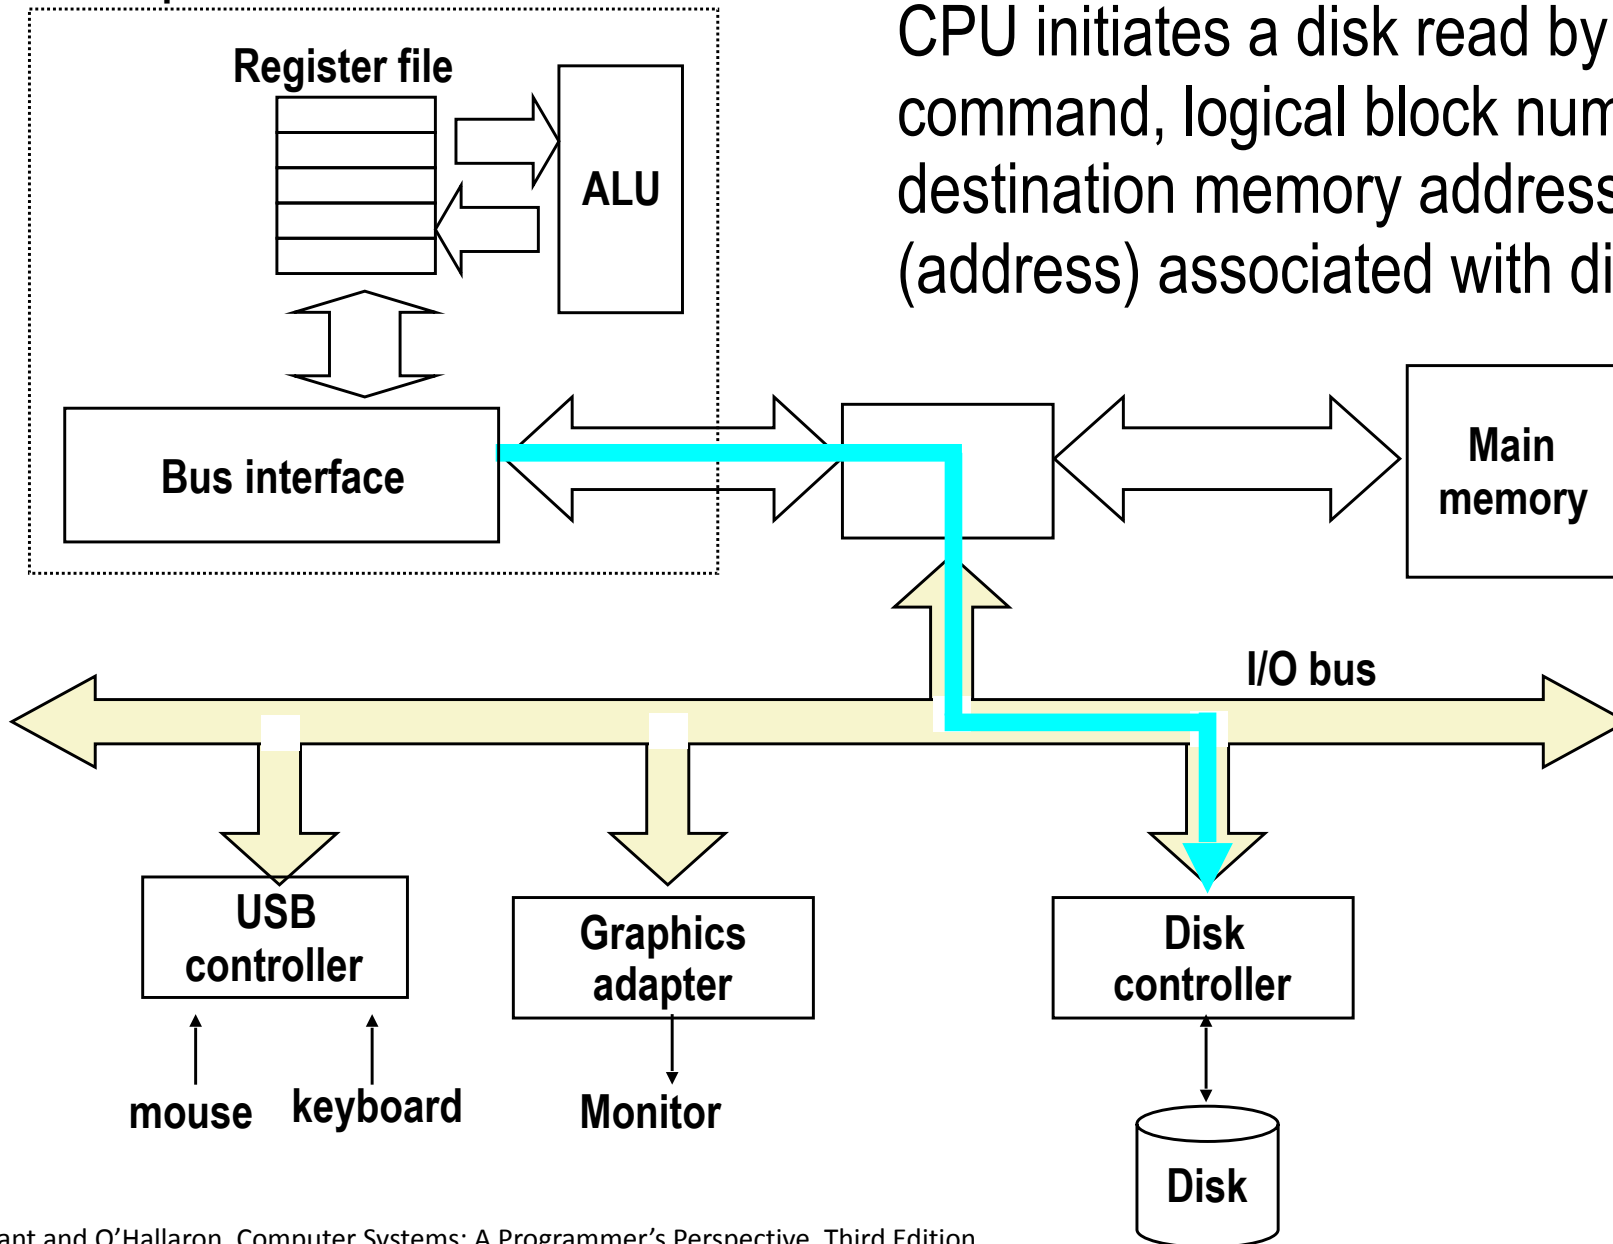
CPU initiates a disk read by writing a command, logical block number, and destination memory address to a port (address) associated with disk controller.

Main memory

I/O bus

USB controller

mouse   keyboard

Graphics adapter

Monitor

Disk controller

Disk

Reading a Disk Sector (2)

CPU chip

Register file

ALU

Disk controller reads the sector and performs a direct memory access (DMA) transfer into main memory.

Bus interface

Main memory

I/O bus

USB controller

Graphics adapter

Disk controller

Mouse **Keyboard**

Monitor
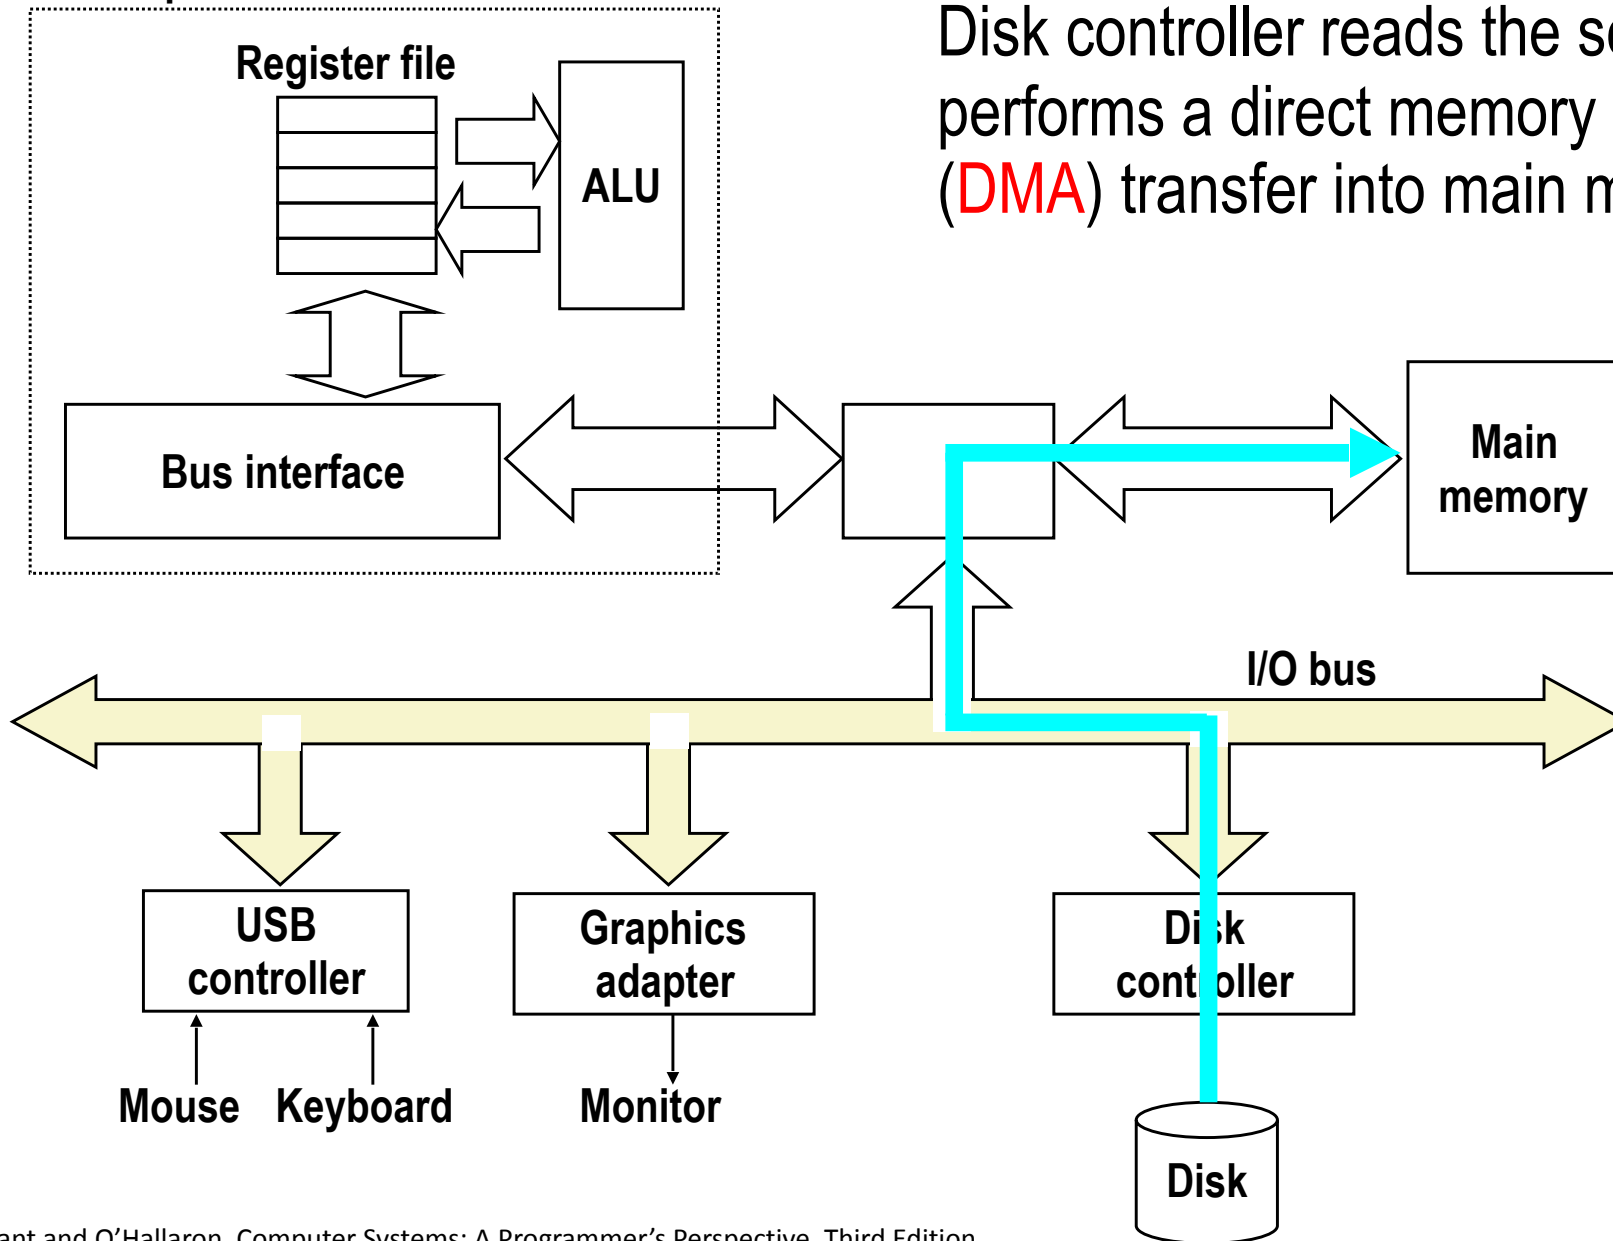
Disk

Reading a Disk Sector (3)

CPU chip

Register file

ALU

Bus interface

When the DMA transfer completes, the disk controller notifies the CPU with an *interrupt* (i.e., asserts a special "interrupt" pin on the CPU)

Main memory

I/O bus

USB controller

Graphics adapter

Disk controller

Mouse   **Keyboard**

Monitor

Disk

Solid State Disks (SSDs)

I/O bus

Requests to read and write logical disk blocks

Solid State Disk (SSD)

Flash
translation layer

Flash memory

Block 0				Block B-1			
Page 0	Page 1	· · ·	Page P-1	· · ·	Page 0	Page 1	· · ·	Page P-1

- **Pages: 512KB to 4KB, Blocks: 32 to 128 pages**

- **Data read/written in units of pages.**

- **Page can be written only after its block has been erased**

- **A block wears out after about 100,000 repeated writes.**

SSD Performance Characteristics

Sequential read tput	550 MB/s	Sequential write tput	470 MB/s
Random read tput	365 MB/s	Random write tput	303 MB/s
Avg seq read time	50 us	Avg seq write time	60 us

- **Sequential access faster than random access**
  - Common theme in the memory hierarchy

- **Random writes are somewhat slower**
  - Erasing a block takes a long time (~1 ms)
  - Modifying a block page requires all other pages to be copied to new block
  - In earlier SSDs, the read/write gap was much larger.

Source: Intel SSD 730 product specification.

SSD Tradeoffs vs Rotating Disks

- **Advantages**
  - No moving parts → faster, less power, more rugged

- **Disadvantages**
  - Have the potential to wear out
    - Mitigated by "wear leveling logic" in flash translation layer
    - E.g. Intel SSD 730 guarantees 128 petabyte (128 x 10^{15} bytes) of writes before they wear out
  - About 6 times more expensive per tera byte (2018)
  - Smaller in size

- **Applications**
  - MP3 players, smart phones, laptops
  - Is moving to desktops and servers

CPU Clock Rates

Inflection point in computer history when designers hit the "Power Wall"

	1985	1990	1995	2003	2005	2010	2017	*2017:1985*
CPU	80286	80386	Pentium	P-4	Core 2	Core i7(n)	Core i7(h)	
Clock rate (MHz)	6	20	150	3,300	2,000	2,500	3,000	500
Cycle time (ns)	166	50	6	0.30	0.50	0.4	0.33	500
Cores	1	1	1	1	2	4	4	4
Effective cycle time (ns)	166	50	6	0.30	0.25	0.10	0.08	2,000

(n) Nehalem processor
(h) Haswell processor

Storage Trends

SRAM

Metric	1985	1990	1995	2000	2005	2010	2015	*2015:1985*
$/MB	2,900	320	256	100	75	60	*320*	*116*
access (ns)	150	35	15	3	2	1.5	*200*	*115*

DRAM

Metric	1985	1990	1995	2000	2005	2010	2015	*2015:1985*
$/MB	880	100	30	1	0.1	0.06	0.02	*44,000*
access (ns)	200	100	70	60	50	40	20	*10*
typical size (MB)	0.256	4	16	64	2,000	8,000	16.000	62,500

Disk

Metric	1985	1990	1995	2000	2005	2010	2015	*2015:1985*
$/GB	100,000	8,000	300	10	5	0.3	0.03	*3,333,333*
access (ms)	75	28	10	8	5	*3*	*3*	25
typical size (GB)	0.01	0.16	1	20	160	1,500	3,000	*300,000*

The CPU-Memory Gap

The gap widens between DRAM, disk, and CPU speeds.

Locality to the Rescue!

The key to bridging this CPU-Memory gap is a fundamental property of computer programs known as locality

Today

- **Performance**
- **Storage technologies and trends**
- **Locality of reference**
- **Caching in the memory hierarchy**

Locality

- **Principle of Locality: Programs tend to use data and instructions with addresses near or equal to those they have used recently**

- **Temporal locality:**
  - Recently referenced items are likely
    to be referenced again in the near future

- **Spatial locality:**
  - Items with nearby addresses tend
    to be referenced close together in time

Locality Example

```
sum = 0;
for (i = 0; i < n; i++)
    sum += a[i];
return sum;
```

- ◾ **Data references**
  - ▪ Reference array elements in succession (stride-1 reference pattern).   **Spatial locality**
  - ▪ Reference variable `sum` each iteration.   **Temporal locality**
- ◾ **Instruction references**
  - ▪ Reference instructions in sequence.   **Spatial locality**
  - ▪ Cycle through loop repeatedly.   **Temporal locality**

Qualitative Estimates of Locality

■ **Claim:** Being able to look at code and get a qualitative sense of its locality is a key skill for a professional programmer.

■ **Question:** Does this function have good locality with respect to array a?

```
int sum_array_rows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
}
```

Locality Example

- **Question:** Does this function have good locality with respect to array a?

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;

}
```

Locality Example

■ **Question: Can you permute the loops so that the function scans the 3-d array** a **with a stride-1 reference pattern (and thus has good spatial locality)?**

```
int sum_array_3d(int a[M][N][N])
{
    int i, j, k, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            for (k = 0; k < N; k++)
                sum += a[k][i][j];
    return sum;
}
```

Memory Hierarchies

- **Some fundamental and enduring properties of hardware and software:**
  - Fast storage technologies cost more per byte, have less capacity, and require more power (heat!).
  - The gap between CPU and main memory speed is widening.
  - Well-written programs tend to exhibit good locality.

- **These fundamental properties complement each other beautifully.**

- **They suggest an approach for organizing memory and storage systems known as a memory hierarchy.**

Today

- **Performance**
- **Storage technologies and trends**
- **Locality of reference**
- **Caching in the memory hierarchy**

Example Memory Hierarchy

Smaller,
faster,
and
costlier
(per byte)
storage
devices

Larger,
slower,
and
cheaper
(per byte)
storage
devices

L0: Regs

CPU registers hold words retrieved from the L1 cache.

L1: L1 cache (SRAM)

L1 cache holds cache lines retrieved from the L2 cache.

L2: L2 cache (SRAM)

L2 cache holds cache lines retrieved from L3 cache

L3: L3 cache (SRAM)

L3 cache holds cache lines retrieved from main memory.

L4: Main memory (DRAM)

Main memory holds disk blocks retrieved from local disks.

L5: Local secondary storage (local disks)

Local disks hold files retrieved from disks on remote servers

L6: Remote secondary storage (e.g., Web servers)

Caches

- *Cache:* **A smaller, faster storage device that acts as a staging area for a subset of the data in a larger, slower device.**

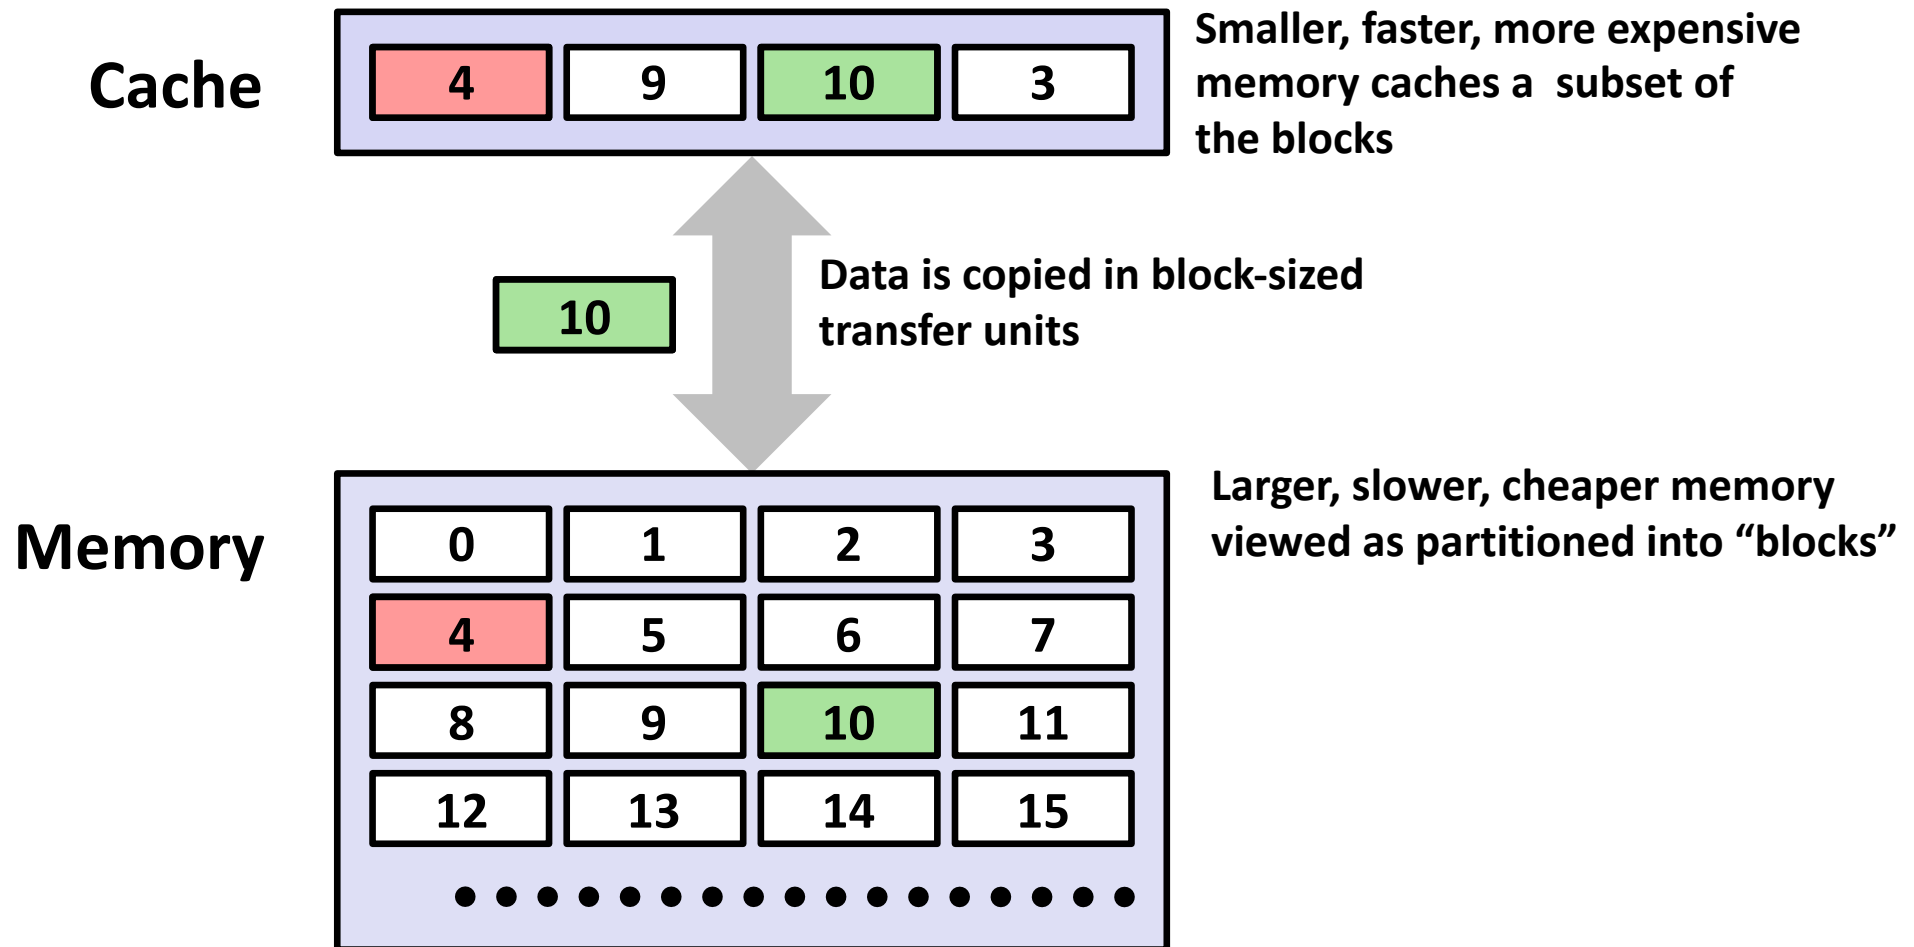- **Fundamental idea of a memory hierarchy:**
  - For each k, the faster, smaller device at level k serves as a cache for the larger, slower device at level k+1.

- **Why do memory hierarchies work?**
  - Because of locality, programs tend to access the data at level k more often than they access the data at level k+1.
  - Thus, the storage at level k+1 can be slower, and thus larger and cheaper per bit.

- *Big Idea:* **The memory hierarchy creates a large pool of storage that costs as much as the cheap storage near the bottom, but that serves data to programs at the rate of the fast storage near the top.**

General Cache Concepts

Cache

| 4 | 9 | 10 | 3 |

Smaller, faster, more expensive memory caches a subset of the blocks

| 10 |

Data is copied in block-sized transfer units

Memory

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

Larger, slower, cheaper memory viewed as partitioned into "blocks"

General Cache Concepts: Hit

Request: 14

Cache

| 8 | 9 | 14 | 3 |

Memory

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

Data in block b is needed

Block b is in cache:
Hit!

General Cache Concepts: Miss

Request: 12

Data in block b is needed

Cache

| 8 | 12 | 14 | 3 |

Block b is not in cache:
Miss!

| 12 |

Request: 12

Block b is fetched from memory

Memory

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

Block b is stored in cache
- Placement policy: determines where b goes
- Replacement policy: determines which block gets evicted (victim)
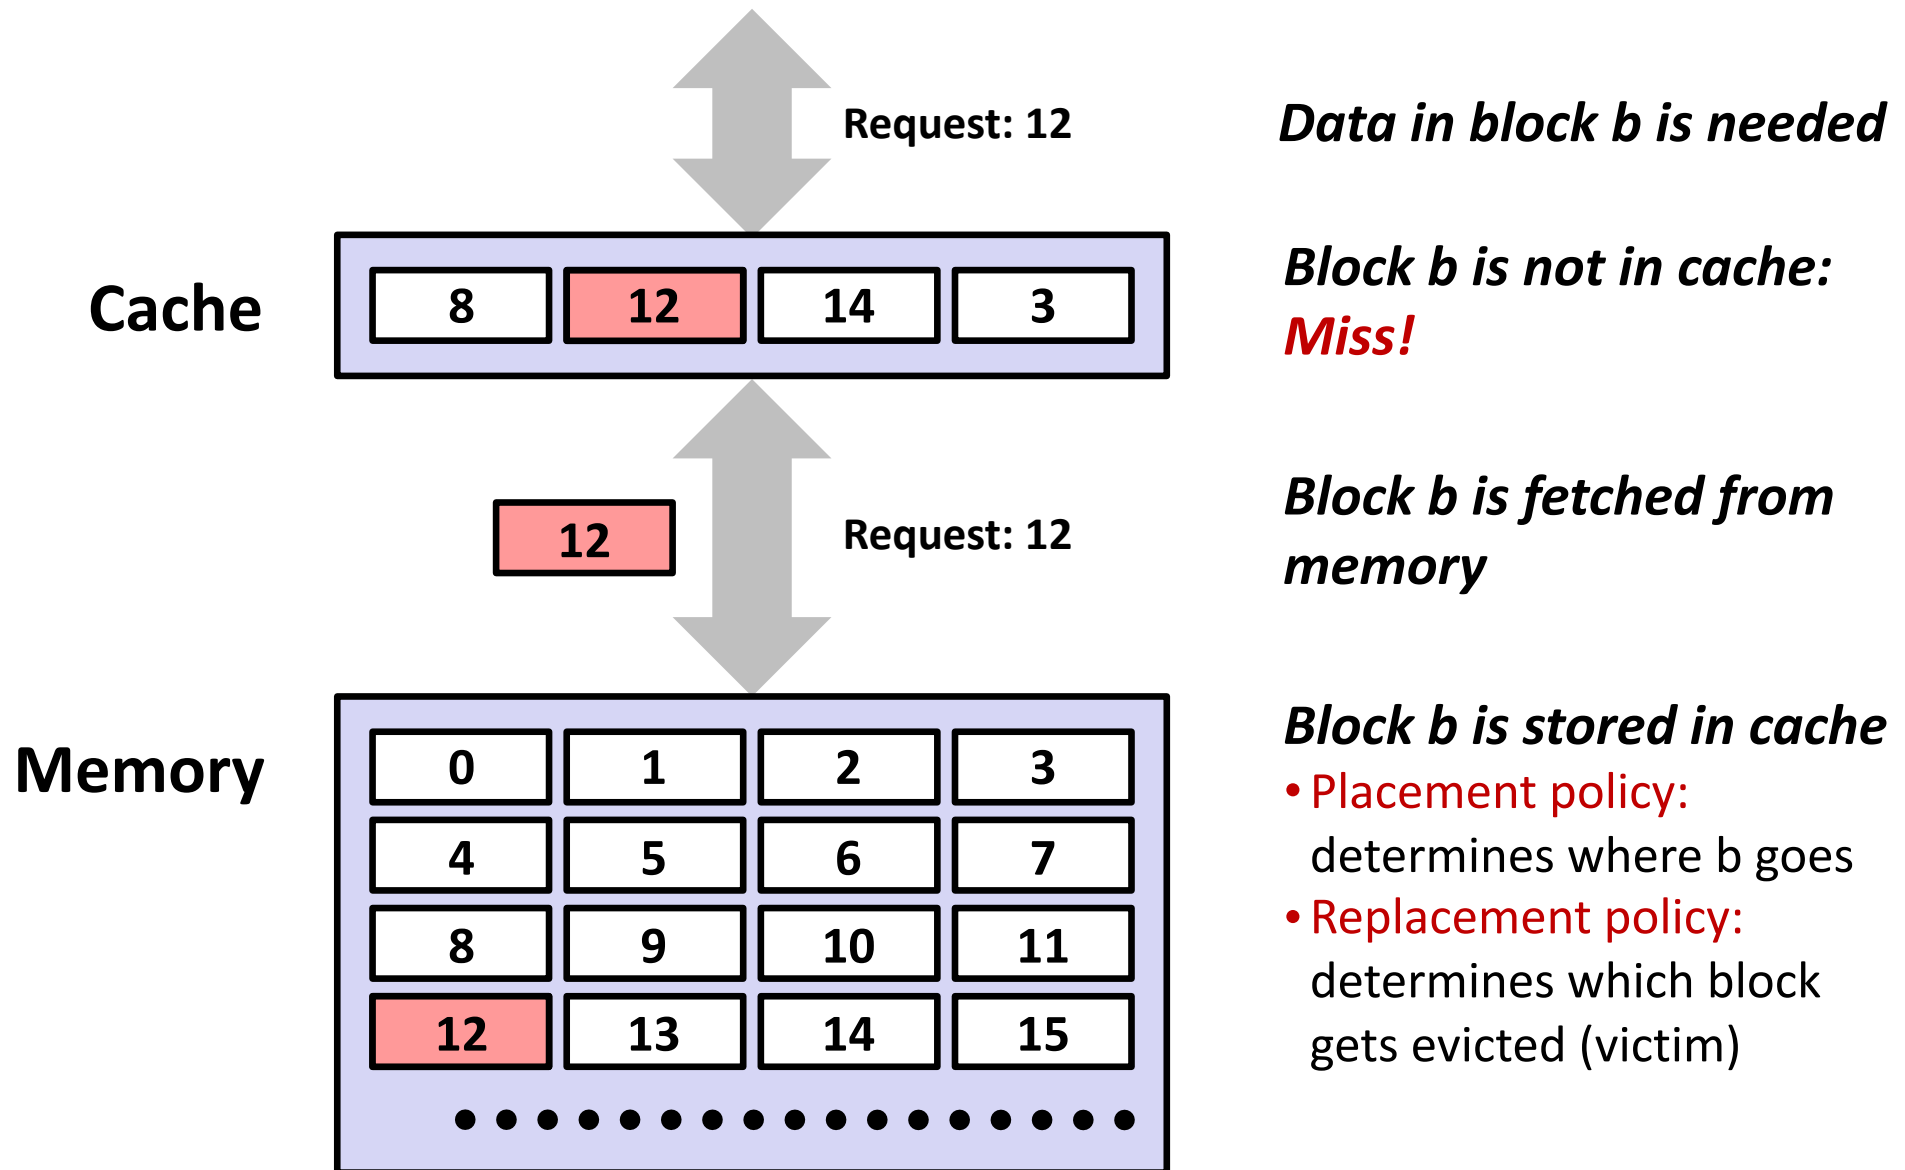
General Caching Concepts:
 Types of Cache Misses

Cold (compulsory) miss

- Cold misses occur because the cache is empty.

Conflict miss

- Most caches limit blocks at level k+1 to a small subset (sometimes a singleton) of the block positions at level k.
  - E.g. Block i at level k+1 must be placed in block (i mod 4) at level k.
- Conflict misses occur when the level k cache is large enough, but multiple data objects all map to the same level k block.
  - E.g. Referencing blocks 0, 8, 0, 8, 0, 8, … would miss every time.

Capacity miss

- Occurs when the set of active cache blocks (working set) is larger than the cache.

Examples of Caching in the Mem. Hierarchy

Cache Type	What is Cached?	Where is it Cached?	Latency (cycles)	Managed By
Registers	4-8 bytes words	CPU core	0	Compiler
TLB	Address translations	On-Chip TLB	0	Hardware MMU
L1 cache	64-byte blocks	On-Chip L1	4	Hardware
L2 cache	64-byte blocks	On-Chip L2	10	Hardware
Virtual Memory	4-KB pages	Main memory	100	Hardware + OS
Buffer cache	Parts of files	Main memory	100	OS
Disk cache	Disk sectors	Disk controller	100,000	Disk firmware
Network buffer cache	Parts of files	Local disk	10,000,000	NFS client
Browser cache	Web pages	Local disk	10,000,000	Web browser
Web cache	Web pages	Remote server disks	1,000,000,000	Web proxy server

Summary

- **The speed gap between CPU, memory and mass storage continues to widen.**

- **Well-written programs exhibit a property called *locality*.**

- **Memory hierarchies based on *caching* close the gap by exploiting locality.**

- **Compilers can help to optimize you code, but**
  - will not change the asymptotic behavior
  - can only give limited help with memory optimizations